

CS 188

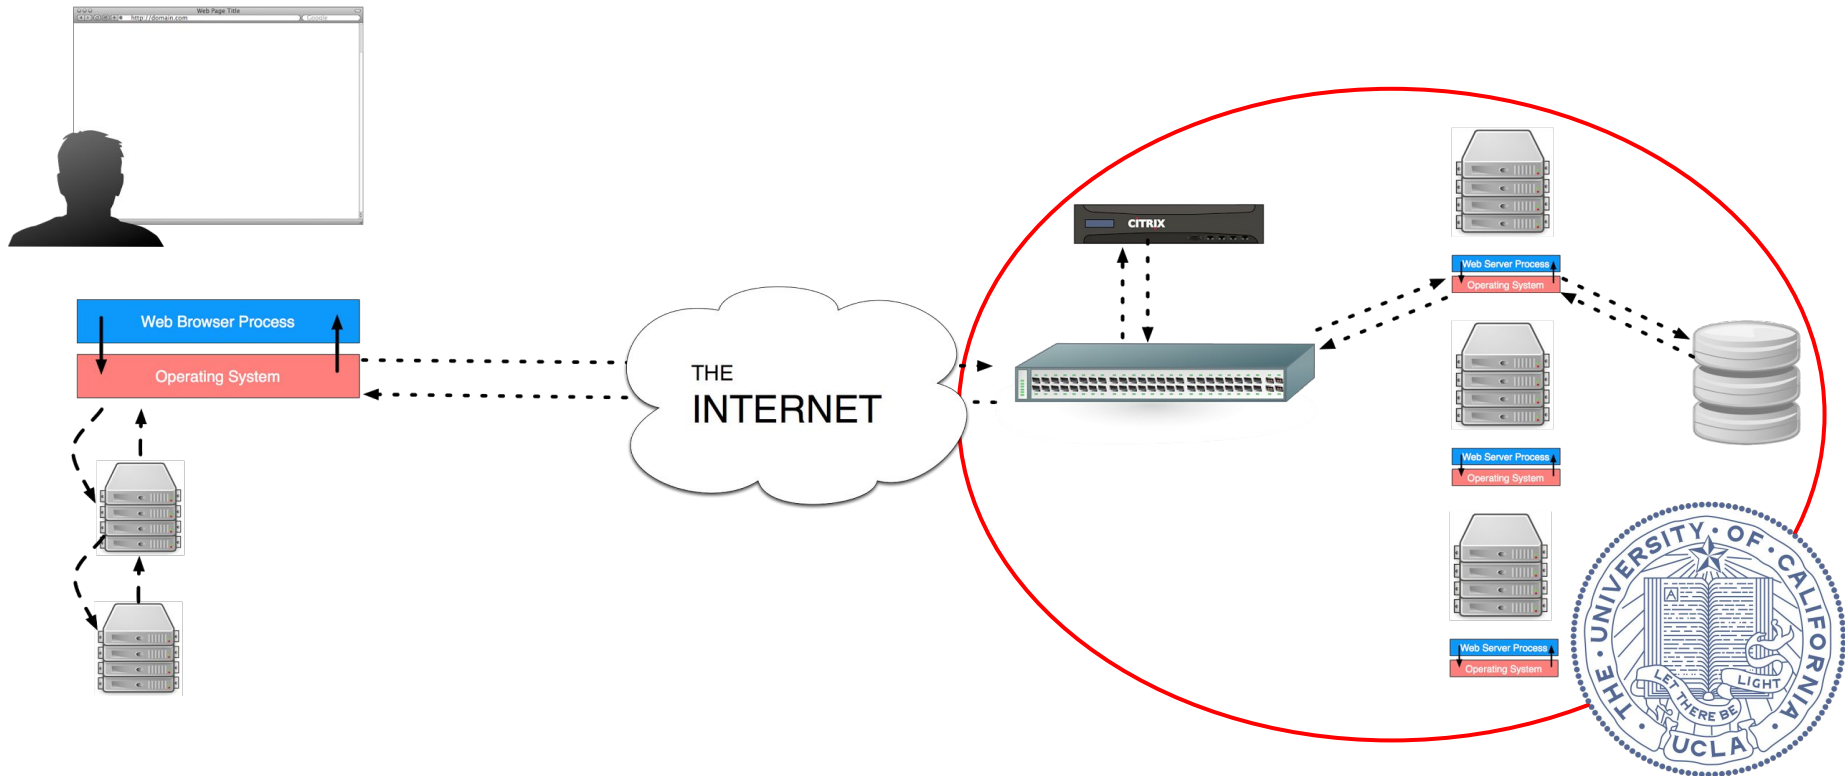
Scalable Internet Services

Andrew Mutz

October 27, 2016



For Today



Motivation

After today you will know how to evaluate the scalability of a deployed application using Tsung.

Today will be interactive, so if you've brought your laptops, please get them out.



Motivation

Let's say we are considering a significant change to our web application.

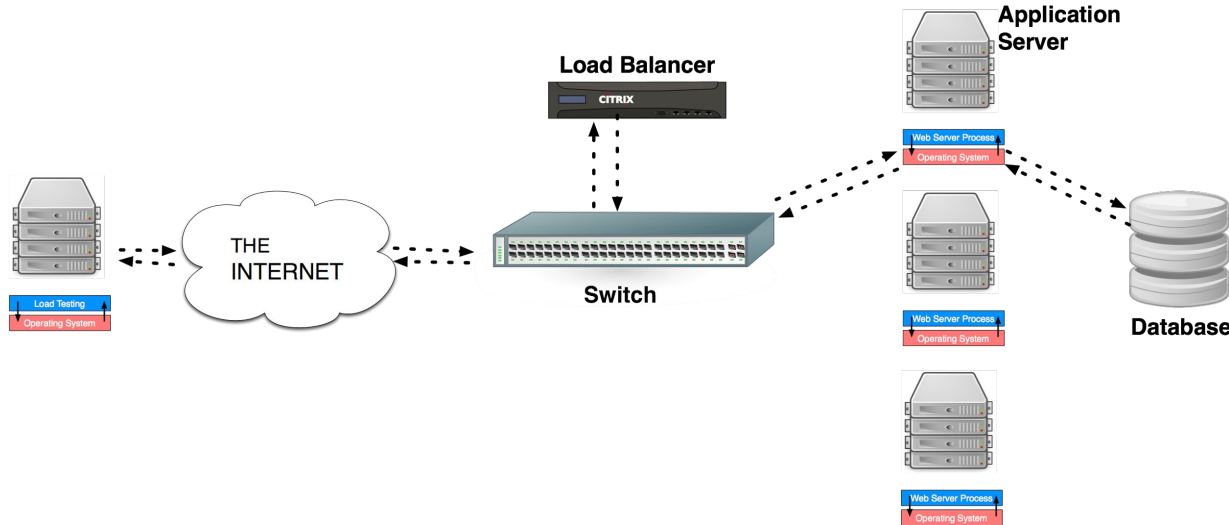
We think it may improve scalability, but our intuition might be wrong.

How should we go about testing this?



Motivation

Let's deploy the system, send actual requests to it, and watch how it responds.



Motivation

What should we observe?

- Response times
- Error rates
- Are our synthetic users able to finish their tasks?



Motivation

Some observations:

- We want a mixture of reads and writes
 - Why is this important?
- Not all users have the same habits
 - Why is this important?
- We want to be able to respond to application output.
 - When is this important?



Motivation

Some load testing tools have high performance

- apachebench, httpperf

Some tools have rich feature sets

- Funkload

Tsung is a good combination of the two



Motivation

We will want one app and testing instance per team. So please:

- Sit with your team
- Deploy one micro instance of your app on EC2
- Deploy one Tsung instance per team
 - I will explain how to do this on the next slide.



Tsung

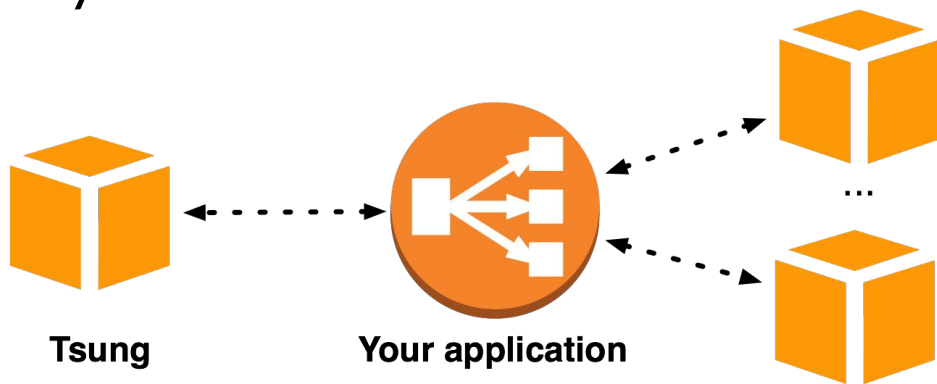
Why Tsung?

- There are many load testing tools out there
- Most make you choose between flexibility and performance
- If you use a low-performance tool, you need to deploy a fleet of machines to do load testing
- Tsung is extremely configurable and delivers high performance



Tsung

- We will do all load testing within AWS.
 - Saves money
 - More predictable
- Don't use T-series (t1 micro, etc) for measurement.
 - Why?



Tsung

Starting up an instance of Tsung testing is easy.

- Create a new CF stack and name it something like TEAM_NAME-tsung
- Use the Tsung template from the utils repo:

Other Templates

- **Tsung:**
This instance provides an installed version of Tsung at your disposal. You will need to copy/rsync over your tsung xml tests.
 - (UCLA) <https://scalableinternetservices.s3.amazonaws.com/Tsung.json>

- Select your TeamName
- Turn off rollback.



Tsung

Once your Tsung stack is up, you will see important information in the “Outputs section”:

Overview	Outputs	Resources	Events	Template	Parameters	Tags	Stack Policy	
Key		Value					Description	
SSH		ssh -i demo.pem ec2-user@ec2-54-200-123-44.us-west-2.amazonaws.com					AppServer SSH connect string	

- SSH tells you how to SSH into your Tsung instance
- The machine address tells you where to go to see results via HTTP



Tsung

Let's SSH into our tsung instance

```
Last login: Tue Oct 27 19:45:28 2015 from cpe-172-250-58-87.socal.res.rr.com
```

```
  _| _|_ )  
 _| ( /  Amazon Linux AMI  
__|\___|__|
```

```
https://aws.amazon.com/amazon-linux-ami/2014.09-release-notes/
```

```
30 package(s) needed for security, out of 137 available
```

```
Run "sudo yum update" to apply all updates.
```

```
Amazon Linux version 2015.09 is available.
```

```
[ec2-user@ip-172-31-27-227 ~]$ ls -l
```

```
total 65624
```

```
drwxrwxr-x  4 ec2-user ec2-user      4096 Oct 26 18:58 opt  
drwxrwxr-x 11 ec2-user ec2-user      4096 Oct 26 18:40 otp_src_R16B03-1  
-rw-rw-r--  1 ec2-user ec2-user 66253556 Jan 24  2014 otp_src_R16B03-1.tar.gz  
drwxr-xr-x  5 root      root        4096 Oct 26 19:02 perl5  
drwxr-xr-x  9 ec2-user ec2-user      4096 Oct 26 18:58 tsung-1.5.0  
-rw-rw-r--  1 ec2-user ec2-user    925026 May 24  2013 tsung-1.5.0.tar.gz  
[ec2-user@ip-172-31-27-227 ~]$
```



Tsung

opt, otp_src_R16B03-1, otp_src_R16B03-1.tar.gz, tsung-1.5.0, tsung-1.5.0.tar.gz can be ignored. They are the installed version of erlang and tsung.

```
Last login: Tue Oct 27 19:45:28 2015 from cpe-172-250-58-87.socal.res.rr.com
```

```
  _| _|_ )  
 _| ( /  Amazon Linux AMI  
__|\___|__|
```

```
https://aws.amazon.com/amazon-linux-ami/2014.09-release-notes/
```

```
30 package(s) needed for security, out of 137 available
```

```
Run "sudo yum update" to apply all updates.
```

```
Amazon Linux version 2015.09 is available.
```

```
[ec2-user@ip-172-31-27-227 ~]$ ls -l
```

```
total 65624
```

drwxrwxr-x	4	ec2-user	ec2-user	4096	Oct 26 18:58	opt
drwxrwxr-x	11	ec2-user	ec2-user	4096	Oct 26 18:40	otp_src_R16B03-1
-rw-rw-r--	1	ec2-user	ec2-user	66253556	Jan 24 2014	otp_src_R16B03-1.tar.gz
drwxr-xr-x	5	root	root	4096	Oct 26 19:02	perl5
drwxr-xr-x	9	ec2-user	ec2-user	4096	Oct 26 18:58	tsung-1.5.0
-rw-rw-r--	1	ec2-user	ec2-user	925026	May 24 2013	tsung-1.5.0.tar.gz

```
[ec2-user@ip-172-31-27-227 ~]$
```



Tsung

Clone your repo to the load testing instance so you can get at your load testing scripts:

```
git clone git@github.com:scalableinternetservices/demo.git
```

Let's kick the tires by testing how www.google.com responds to light load:

```
tsung -n -f test.xml start
```

```
[ec2-user@ip-172-31-27-227 ~]$ tsung -f test.xml start
Starting Tsung
"Log directory is: /home/ec2-user/.tsung/log/20151027-1955"
[ec2-user@ip-172-31-27-227 ~]$
```

NOTICE: please add -n flag when running tsung



Tsung

This will take a little time to run. Afterwards we compile the report:

```
cd /home/ec2-user/.tsung/log/20150504-0444; tsung_stats.pl
```

```
[ec2-user@ip-172-31-27-227 ~]$ tsung -f test.xml start
Starting Tsung
"Log directory is: /home/ec2-user/.tsung/log/20151027-1955"
[ec2-user@ip-172-31-27-227 ~]$ cd /home/ec2-user/.tsung/log/20151027-1955
[ec2-user@ip-172-31-27-227 20151027-1955]$ tsung_stats.pl
creating subdirectory data
creating subdirectory gnuplot_scripts
creating subdirectory images
warn, last interval (6) not equal to the first, use the first one (10)
No data for Bosh
No data for Match
No data for Event
No data for Async
No data for Errors
[ec2-user@ip-172-31-27-227 20151027-1955]$
```



Tsung

Once we have compiled the report, switch over to the web interface:

Overview			Outputs			Resources			Events			Template			Parameters			Tags			Stack Policy		
Key			Value			Description																	
SSH			ssh -i demo.pem ec2-user@ec2-54-200-123-44.us-west-2.com pute.amazonaws.com			AppServer SSH connect string																	

This just serves up the files from ~/.tsung/log/ through a web interface



Tsung

Once we have compiled the report, switch over to the web interface:

Index of /

<u>Name</u>	<u>Last modified</u>
Parent Directory	2015/05/04 02:14
20150504-0444/	2015/05/04 04:47

*WEBrick/1.3.1 (Ruby/2.1.5/2014-11-13)
at ec2-52-24-67-201.us-west-2.compute.amazonaws.com:80*



Tsung

Once we have compiled the report, switch over to the web interface:

Index of /20150504-0444/

<u>Name</u>	<u>Last modified</u>
Parent Directory	2015/05/04 04:44
data/	2015/05/04 04:47
gnuplot.log	2015/05/04 04:47
gnuplot scripts/	2015/05/04 04:47
graph.html	2015/05/04 04:47
images/	2015/05/04 04:47
match.log	2015/05/04 04:44
report.html	2015/05/04 04:47
test.xml	2015/05/04 04:44
tsung.log	2015/05/04 04:45
tsung_controller@ip-172..	2015/05/04 04:45

WEBrick/1.3.1 (Ruby/2.1.5/2014-11-13)
at ec2-52-24-67-201.us-west-2.compute.amazonaws.com:80



Tsung

Once we have compiled the report, switch over to the web interface:

Tsung
version 1.5.0

[Stats Report](#)

- [Main statistics](#)
- [Transactions](#)
- [Network Throughput](#)
- [Counters](#)
- [HTTP status](#)

[Graphs Report](#)

- [Response times](#)
- [Throughput graphs](#)
- [Simultaneous Users](#)
- [HTTP status](#)

[XML Config file](#)

tsung - Statistics

Main Statistics

Name	highest 10sec mean	lowest 10sec mean	Highest Rate	Mean	Count
connect	0.42 sec	9.26 msec	0.5 / sec	0.35 sec	6
page	0.57 sec	66.26 msec	0.4 / sec	0.40 sec	6
request	0.57 sec	66.26 msec	0.4 / sec	0.40 sec	6
session	0.48 sec	68.15 msec	0.5 / sec	0.41 sec	6

Transactions Statistics

Name	highest 10sec mean	lowest 10sec mean	Highest Rate	Mean	Count
------	--------------------	-------------------	--------------	------	-------

Network Throughput

Name	Highest Rate	Total
size_rcv	68.01 Kbits/sec	127.47 KB
size_sent	560 bits/sec	840 B

Counters Statistics

Name	Highest Rate	Total number
------	--------------	--------------

Name	Max
connected	1
finish_users_count	6
users	2
users_count	6

HTTP return code

Code	Highest Rate	Total number
200	0.4 / sec	6

We will go over this report in more detail later today



Tsung - test.xml

Let's go over the basic xml file:

```
<?xml version="1.0"?><tsung loglevel="notice" version="1.0">
  <clients>
    <client host="localhost" use_controller_vm="true" maxusers="15000"/>
  </clients>
  <servers>
    <server host="www.google.com" port="80" type="tcp"/>
  </servers>
  <load>
    <arrivalphase phase="1" duration="10" unit="second">
      <users arrivalrate="1" unit="second"/>
    </arrivalphase>
  </load>
```



Tsung - test.xml

XML Boilerplate

```
<?xml version="1.0"?><tsung loglevel="notice" version="1.0">
  <clients>
    <client host="localhost" use_controller_vm="true" maxusers="15000"/>
  </clients>
  <servers>
    <server host="www.google.com" port="80" type="tcp"/>
  </servers>
  <load>
    <arrivalphase phase="1" duration="10" unit="second">
      <users arrivalrate="1" unit="second"/>
    </arrivalphase>
  </load>
```



Tsung - test.xml

Client-side configuration. Maxusers is the maximum number of simulated users.

```
<?xml version="1.0"?><tsung loglevel="notice" version="1.0">  
  <clients>  
    <client host="localhost" use_controller_vm="true" maxusers="15000"/>  
  </clients>  
  <servers>  
    <server host="www.google.com" port="80" type="tcp"/>  
  </servers>  
  <load>  
    <arrivalphase phase="1" duration="10" unit="second">  
      <users arrivalrate="1" unit="second"/>  
    </arrivalphase>  
  </load>
```



Tsung - test.xml

Server configuration: where we are directing load.

```
<?xml version="1.0"?><tsung loglevel="notice" version="1.0">
  <clients>
    <client host="localhost" use_controller_vm="true" maxusers="15000"/>
  </clients>
  <servers>
    <server host="www.google.com" port="80" type="tcp"/>
  </servers>
  <load>
    <arrivalphase phase="1" duration="10" unit="second">
      <users arrivalrate="1" unit="second"/>
    </arrivalphase>
  </load>
```



Tsung - test.xml

Defining phases of user arrival rates

```
<?xml version="1.0"?><tsung loglevel="notice" version="1.0">
  <clients>
    <client host="localhost" use_controller_vm="true" maxusers="15000"/>
  </clients>
  <servers>
    <server host="www.google.com" port="80" type="tcp"/>
  </servers>
  <load>
    <arrivalphase phase="1" duration="10" unit="second">
      <users arrivalrate="1" unit="second"/>
    </arrivalphase>
  </load>
```



Tsung - test.xml

A section to set options (timeouts, useragents)

```
<options>
```

```
  <option name="glocal_ack_timeout" value="2000"/>
```

```
  <option type="ts_http" name="user_agent">
```

```
    <user_agent probability="100">Mozilla/5.0 (Windows; U; Windows NT 5.2; fr-FR; rv:1.7.8) Gecko/20050511  
    Firefox/1.0.4</user_agent>
```

```
  </option>
```

```
</options>
```

```
<sessions>
```

```
  <session name="http-example" probability="100" type="ts_http">
```

```
    <request>
```

```
      <http url="/" version="1.1" method="GET"/>
```

```
    </request>
```

```
  </session>
```

```
</sessions>
```

```
</tsung>
```



Tsung - test.xml

We define the actual series of requests that a user will perform.

```
<options>
  <option name="glocal_ack_timeout" value="2000"/>
  <option type="ts_http" name="user_agent">
    <user_agent probability="100">Mozilla/5.0 (Windows; U; Windows NT 5.2; fr-FR; rv:1.7.8) Gecko/20050511
    Firefox/1.0.4</user_agent>
  </option>
</options>
<sessions>
  <session name="http-example" probability="100" type="ts_http">
    <request>
      <http url="/" version="1.1" method="GET"/>
    </request>
  </session>
</sessions>
</tsung>
```



Tsung - test.xml

Lets change our tests to point to our server

```
<?xml version="1.0"?><tsung loglevel="notice" version="1.0">
  <clients>
    <client host="localhost" use_controller_vm="true" maxusers="15000"/>
  </clients>
  <servers>
    <server host="ec2-52-24-120-129.us-west-2.compute.amazonaws.com" port="80"
type="tcp"/>
  </servers>
  <load>
    <arrivalphase phase="1" duration="10" unit="second">
      <users arrivalrate="1" unit="second"/>
    </arrivalphase>
  </load>
```



Tsung - test.xml

```
<load>
  <arrivalphase phase="1" duration="30" unit="second">
    <users arrivalrate="1" unit="second"></users>
  </arrivalphase>
  <arrivalphase phase="2" duration="30" unit="second">
    <users arrivalrate="2" unit="second"></users>
  </arrivalphase>
  <arrivalphase phase="3" duration="30" unit="second">
    <users arrivalrate="4" unit="second"></users>
  </arrivalphase>
  <arrivalphase phase="4" duration="30" unit="second">
    <users arrivalrate="8" unit="second"></users>
  </arrivalphase>
</load>
```

For examples of more complex load testing, see the demo app's `load_tests/critical.xml`

Increasing the rate of user creation over time.



Tsung - test.xml

```
<load>
  <arrivalphase phase="1" duration="30" unit="second">
    <users arrivalrate="1" unit="second"></users>
  </arrivalphase>
  <arrivalphase phase="2" duration="30" unit="second">
    <users arrivalrate="2" unit="second"></users>
  </arrivalphase>
  <arrivalphase phase="3" duration="30" unit="second">
    <users arrivalrate="4" unit="second"></users>
  </arrivalphase>
  <arrivalphase phase="4" duration="30" unit="second">
    <users arrivalrate="8" unit="second"></users>
  </arrivalphase>
</load>
```

For examples of more complex load testing, see the demo app's `load_tests/critical.xml`

Increasing the rate of user creation over time.



Tsung - test.xml

```
<sessions>
  <session name="http-example" probability="100" type="ts_http">

    <!-- start out at the dashboard. -->
    <request>
      <http url="/" version='1.1' method='GET'></http>
    </request>
```

For examples of more complex load testing, see the demo app's `load_tests/critical.xml`

The session defines the virtual user that you will be simulating.

This one starts at “/”



Tsung - test.xml

```
<!-- wait for up to 2 seconds, user is looking at posts -->  
<thinktime value="2" random="true"></thinktime>
```

For examples of more complex load testing, see the demo app's `load_tests/critical.xml`

Insert realistic random wait times in your simulations.



Tsung - test.xml

```
<!-- create a random number to make a unique community name -->
<setdynvars sourcetype="random_number" start="1000"
end="999999">
  <var name="random_community_name" />
</setdynvars>

<!-- post to /communities to create a community.
remember the url of the created community so we can delete
it later -->
<request subst="true">
  <http
    url='/communities'
    version='1.1'
    method='POST'
    contents='community%5Bname%5D=community_name_%_random_community_name%_
&amp;commit=Create+Community'>
  </http>
</request>
```

For examples of more complex load testing, see the demo app's [load_tests/critical.xml](#)

Deal with uniqueness constraints by defining dynamic variables.

Insert dynamic variables by using %% syntax.



Tsung - test.xml

```
<request subst="true">
  <dyn_variable name="created_submission_url" re="[L1]ocation:
(http://.*)\r"/>
  <dyn_variable name="created_submission_id" re="[L1]ocation:
http://.*/submissions/(.*)\r"/>
  <http
    url='/submissions'
    version='1.1'
    method='POST'

    contents='submission%5Btitle%5D=link_%_random_submission_name%&su
bmission%5Burl%5D=http%3A%2F%2Fwww.article.com%2F%_random_submission_n
ame%&submission%5Bcommunity_id%5D=%_%created_community_id%&co
mmit=Create+Submission'

  </http>
</request>
```

For examples of more complex load testing, see the demo app's `load_tests/critical.xml`

At times, you will need the output of one request to feed into the next.

Use dynamic variables for this



Tsung - test.xml

```
<!-- Uncomment the following to debug print your dynamic variables
-->

<!--
<setdynvars sourcetype="eval" code="fun( {Pid, DynVars} ) ->
  io:format([126, $p, 126, $n, 126, $n], [DynVars]),
  ok end.">
  <var name="dump" />
</setdynvars>
-->
```

For examples of more complex load testing, see the demo app's `load_tests/critical.xml`

If you have difficulty with your dynamic variables, use this code to debug



Tsung - test.xml

Name

submissions

51

application-13e64a74b7ab4dfdd9d96f8038945dd.css

application-f221e1ec5ab975eeadb83b9e995b0d3.js

bootstrap.min.css

bootstrap-theme.min.css

bootstrap.min.js

nr-632.min.js

33e249c12f7a=5272963&pl=1430718915465&v=632.2b17625&to=JQ5dQUUMWwIVR...

favicon.ico

Headers Preview Response Cookies Timing

General

Remote Address: 52.24.120.129:80

Request URL: http://ec2-52-24-120-129.us-west-2.compute.amazonaws.com/submissions

Request Method: POST

Status Code: 302 Found

Response Headers view source

Cache-Control: no-cache

Connection: keep-alive

Content-Length: 137

Content-Type: text/html; charset=utf-8

Date: Mon, 04 May 2015 05:55:15 GMT

Location: http://ec2-52-24-120-129.us-west-2.compute.amazonaws.com/submissions/51

Server: WEBrick/1.3.1 (Ruby/2.1.5/2014-11-13)

Set-Cookie: _demo_session=QULyenRocLVdC8lzb1hNMwRwXZNVd3008GdXZ5bEvNzGjhb0hjhZ4VepITy9hNG0wb0M2K24yUmp0dn1WeEp0HFKd5LCTFm5nLWGF0TVVYVjZ2aFfc1B2NDJJMEJzanlctYUJ3TDcyd090RwX0d0W0raIordnZnbEvU5Igs5LR1WVJH0HwCv1JWHL1Mw4yU2NwHm5Dh0XxwU5Wk53WwTbPpNX0E0wFRWVpLd3p1LzVSUBw1J61nQ3RJ3Hd9Qk3wac92WmLeF72hzRJY1cxdlFwc0ndwhURE0XvAV2HCTEBzeFmFmBvOmPwPNV2VXbmZuZUkuZzdRdphWkxkZwFjWV12y9CSUxaZwC9P58tzc2PaHoyV3c4d3FCamul2dheWZadz09--3ec242e7eb210b1c8177fb2da1b421138de34146; path=/; HttpOnly

Set-Cookie: request_method=POST; path=/

X-Content-Type-Options: nosniff

X-Frame-Options: SAMEORIGIN

X-Request-Id: 65906fbc-b0f7-4587-892c-1e8e7e7c81b4

X-Runtime: 0.029938

X-Xss-Protection: 1; mode=block

Request Headers view source

Accept: text/html,application/xhtml+xml,application/xml;q=0.9,image/webp,*/*;q=0.8

Accept-Encoding: gzip, deflate

Accept-Language: en-US,en;q=0.8

Cache-Control: max-age=0

Connection: keep-alive

Content-Length: 274

Content-Type: application/x-www-form-urlencoded

Cookie: _demo_session=Mndyb3F1ehJad0NjWGDMEF3U5XVkrRPFVHKZ4b0Rg0E84eUHQjVjM1BHa2RuRVgxlEvlcTNJV2xLS1hZNR0TU0wQ1dqETfYkKhEdzFUS2M5JZYzAvkZolzRXb18ySmIjce9Y0W1Hb09UmwCM3VRTG10b0ZB0EFmTE9NZRXvklZn0pYm14amVeFyVehQ3PT0LXPZ5Z0zNEtKdXRG5lpIdESXOG5TR39PQ303D0--61704676792975833c05bed79338dbde59645a3

Host: ec2-52-24-120-129.us-west-2.compute.amazonaws.com

For examples of more complex load testing, see the demo app's [load_tests/critical.xml](#)

If you are unsure how to simulate your application, use the browser to get a firm idea of the exact HTTP requests.



Form Data view source view decoded

utf8: %E2%9C%93

authenticity_token: yQ87%2F0xfj%2BkNMJ5Lp6r0hs60BeEDI1FY25HeJrf0jL9j7QdVuSRPCJ8Q0Ii8ZD8x%2Fe46eIS9cZkQ0GHFGzQQQ%3D%3D

submission%5Btitle%5D: this+is+a+test

submission%5Burl%5D: http%3A%2F%2Fwww.testing.com

submission%5Bcommunity_id%5D: 1

commit: Create+Submission

Tsung - test.xml

At any point in time you can see how tsung is doing, by typing “tsung status” from another terminal

```
[ec2-user@ip-172-31-27-227 ~]$ tsung status
Tsung is running [OK]
Current request rate: 14.76 req/sec
Current users: 28
Current connected users: 29
Current phase: 3
```



Tsung - Output

Main Statistics

Name	highest 10sec mean	lowest 10sec mean	Highest Rate	Mean	Count
connect	0.42 sec	9.26 msec	0.5 / sec	0.35 sec	6
page	0.57 sec	66.26 msec	0.4 / sec	0.40 sec	6
request	0.57 sec	66.26 msec	0.4 / sec	0.40 sec	6
session	0.48 sec	68.15 msec	0.5 / sec	0.41 sec	6

- **request** Response time for each request.
- **page** Response time for each set of requests (a page is a group of request not separated by a thinktime).
- **connect** Duration of the connection establishment.
- **session** Duration of a user's session.



Tsung - Output

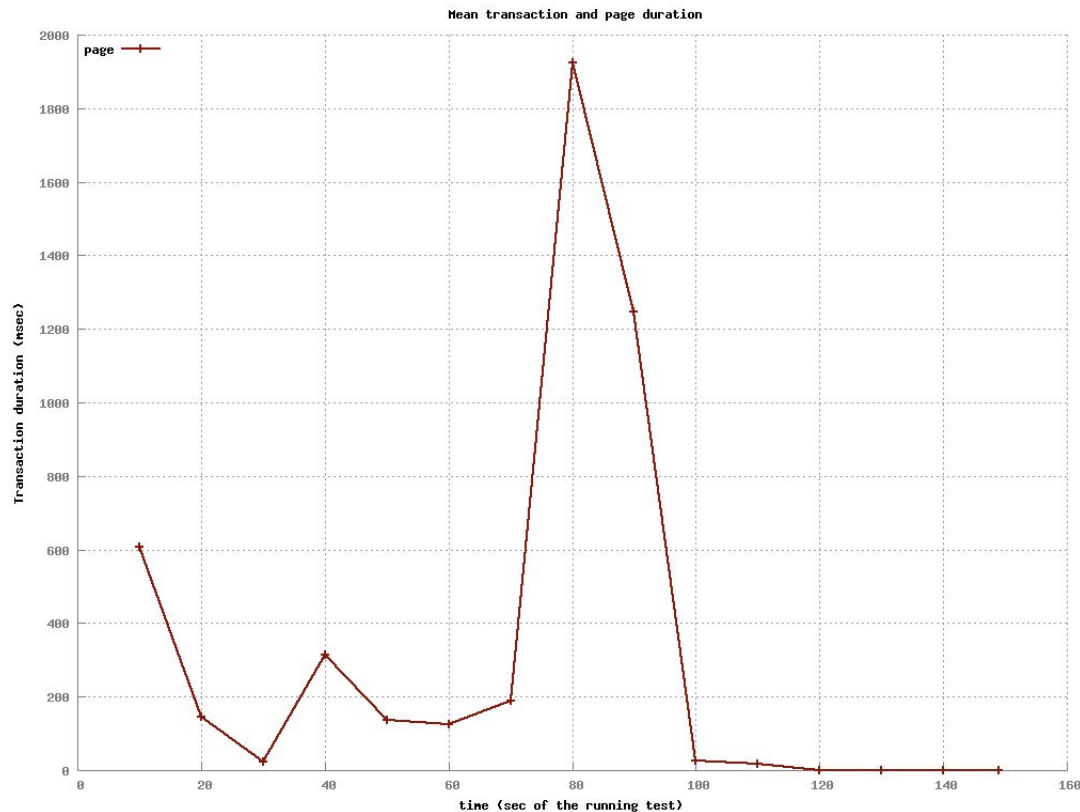
HTTP return code

Code	Highest Rate	Total number
200	0.4 / sec	6

- Make sure you are getting back good status codes.
 - 200s and 300s are good
 - 400s and 500s are usually bad



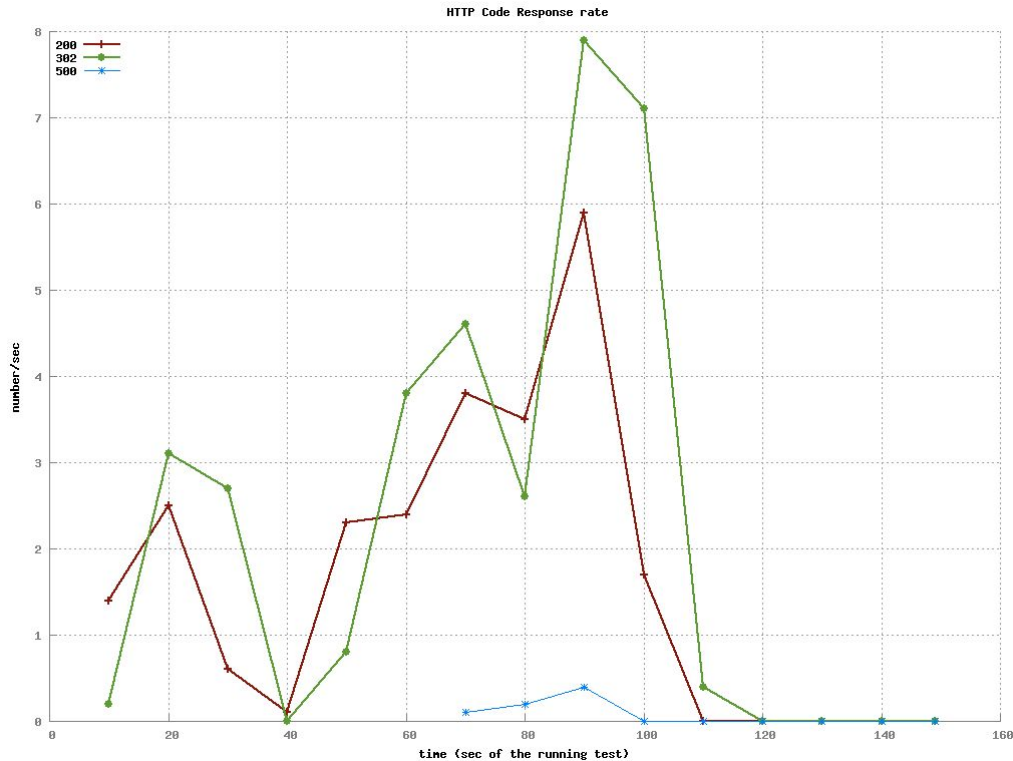
Tsung - Output



- Response times



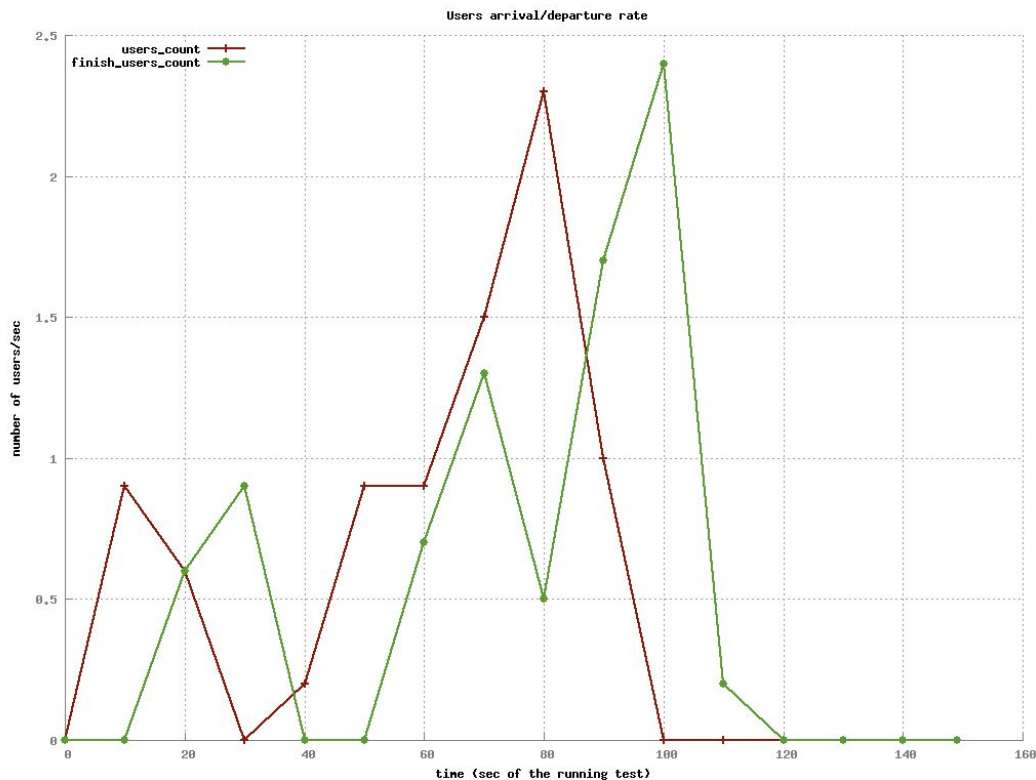
Tsung - Output



- Error rates



Tsung - Output



- Are our synthetic users able to finish their tasks?



Tsung

Always remember that your Tsung instances will go down automatically (8 hours), so please scp off any important data or results immediately.



For Next Time...

Attempt to create a simple load testing script for your current app

Keep completing stories!

